

Title : "Quantum Virtual Internship - Retail Strategy and Analytics - Task 1"

Load required Libraries.

```
In [84]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import re
from collections import Counter
```

Import Dataset

```
In [20]: # Install dependencies as needed:
# pip install kagglehub[pandas-datasets]
import kagglehub
from kagglehub import KaggleDatasetAdapter

# Set the path to the file you'd like to load
file_path = "QVI_transaction_data.csv"

# Load the latest version
transaction_data = kagglehub.dataset_load(
    KaggleDatasetAdapter.PANDAS,
    "preethis14/qvi-transaction-data",
    file_path,
    # Provide any additional arguments like
    # sql_query or pandas_kwargs. See the
    # documentation for more information:
    # https://github.com/Kaggle/kagglehub/blob/main/README.md#kaggledatasetadapterpandas
)

print("First 5 records:", transaction_data.head())
```

		DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	PROD_NBR	\
0	43390	1	1000	1	5		
1	43599	1	1307	348	66		
2	43605	1	1343	383	61		
3	43329	2	2373	974	69		
4	43330	2	2426	1038	108		

			PROD_NAME	PROD_QTY	TOT_SALES
0	Natural Chip	Compny SeaSalt	175g	2	6.0
1		CCs Nacho Cheese	175g	3	6.3
2	Smiths Crinkle Cut	Chips Chicken	170g	2	2.9
3	Smiths Chip Thinly	S/Cream&Onion	175g	5	15.0
4	Kettle Tortilla	ChpsHny&Jlpno	Chili 150g	3	13.8

```
In [22]: # Install dependencies as needed:
# pip install kagglehub[pandas-datasets]
import kagglehub
from kagglehub import KaggleDatasetAdapter

# Set the path to the file you'd like to load
file_path = "QVI_purchase_behaviour.csv"

# Load the latest version
customer_data = kagglehub.dataset_load(
    KaggleDatasetAdapter.PANDAS,
    "preethis14/qvi-purchase-behaviour",
    file_path,
    # Provide any additional arguments like
    # sql_query or pandas_kwargs. See the
    # documentation for more information:
    # https://github.com/Kaggle/kagglehub/blob/main/README.md#kaggledatasetadapterpandas
)

print("First 5 records:", customer_data.head())
```

		LYLTY_CARD_NBR	LIFESTAGE	PREMIUM_CUSTOMER
0	1000	YOUNG SINGLES/COUPLES		Premium
1	1002	YOUNG SINGLES/COUPLES		Mainstream
2	1003	YOUNG FAMILIES		Budget
3	1004	OLDER SINGLES/COUPLES		Mainstream
4	1005	MIDAGE SINGLES/COUPLES		Mainstream

Examine the transaction dataset. Look for the format of each column.

```
In [28]: data_types = transaction_data.dtypes
print(data_types)
```

```

DATE                int64
STORE_NBR           int64
LYLTY_CARD_NBR      int64
TXN_ID              int64
PROD_NBR            int64
PROD_NAME           object
PROD_QTY            int64
TOT_SALES           float64
dtype: object

```

The dates are in Excel serial date format. Conversion to the required format.

```

In [63]: excel_dates = transaction_data['DATE']

dates = pd.to_datetime(excel_dates, origin='1899-12-30', unit='D')

transaction_data['DATE'] = dates

print(dates)

```

```

0      2018-10-17
1      2019-05-14
2      2019-05-20
3      2018-08-17
4      2018-08-18
...
264831 2019-03-09
264832 2018-08-13
264833 2018-11-06
264834 2018-12-27
264835 2018-09-22
Name: DATE, Length: 246742, dtype: datetime64[ns]

```

Summary of product name. examine the product name column.

```

In [40]: product_name = transaction_data['PROD_NAME']
print(product_name.info())
print(product_name.describe(include = "all"))

```

```

<class 'pandas.core.series.Series'>
RangeIndex: 264836 entries, 0 to 264835
Series name: PROD_NAME
Non-Null Count  Dtype
-----
264836 non-null object
dtypes: object(1)
memory usage: 2.0+ MB
None
count                264836
unique                 114
top      Kettle Mozzarella  Basil & Pesto 175g
freq                3304
Name: PROD_NAME, dtype: object

```

Checking for any inconsistencies in product name column.

```

In [51]: def extract_clean_words(name):
          cleaned_name = re.sub(r'^a-zA-Z\s', '', name)
          words = cleaned_name.split()
          return ' '.join(words)

transaction_data['CLEAN_PROD_WORDS'] = transaction_data['PROD_NAME'].apply(extract_clean_words)

print(transaction_data[['PROD_NAME', 'CLEAN_PROD_WORDS']].head())

```

```

          PROD_NAME \
0    Natural Chip      Compny SeaSalt175g
1              CCs Nacho Cheese  175g
2    Smiths Crinkle Cut Chips Chicken 170g
3    Smiths Chip Thinly S/Cream&Onion 175g
4    Kettle Tortilla ChpsHnyJlpno Chili 150g

          CLEAN_PROD_WORDS
0    Natural Chip Compny SeaSaltg
1              CCs Nacho Cheese g
2    Smiths Crinkle Cut Chips Chicken g
3    Smiths Chip Thinly SCreamOnion g
4    Kettle Tortilla ChpsHnyJlpno Chili g

```

Most common words by counting the number of times a word appears and sorting them by this frequency in order of highest to lowest frequency

```

In [55]: all_words = ' '.join(transaction_data['CLEAN_PROD_WORDS']).split()

word_counts = Counter(all_words)

def get_word_frequency(name):

```

```

words = name.split()
return sorted([word_counts[word] for word in words], reverse = True)

transaction_data['WORD_FREQUENCY'] = transaction_data['CLEAN_PROD_WORDS'].apply(get_word_frequency)

print(transaction_data[['PROD_NAME', 'CLEAN_PROD_WORDS', 'WORD_FREQUENCY']].head())

```

```

      PROD_NAME \
0  Natural Chip    Compny SeaSalt175g
1      CCs Nacho Cheese    175g
2  Smiths Crinkle Cut Chips Chicken 170g
3  Smiths Chip Thinly S/Cream&Onion 175g
4  Kettle Tortilla ChpsHny&Jlpno Chili 150g

      CLEAN_PROD_WORDS \
0  Natural Chip Compny SeaSaltg
1      CCs Nacho Cheese g
2  Smiths Crinkle Cut Chips Chicken g
3  Smiths Chip Thinly SCreamOnion g
4  Kettle Tortilla ChpsHnyJlpno Chili g

      WORD_FREQUENCY
0      [18645, 6050, 1468, 1468]
1      [246628, 27890, 4658, 4551]
2  [246628, 49770, 28860, 23960, 20754, 15407]
3      [246628, 28860, 18645, 7507, 1473]
4      [246628, 41288, 9580, 3296, 3296]

```

Checking for chips and not chips

```

In [56]: transaction_data['SALSA'] = transaction_data['PROD_NAME'].str.contains('salsa', case=False, na=False)

transaction_data = transaction_data[transaction_data['SALSA'] == False]

transaction_data = transaction_data.drop(columns=['SALSA'])

print(transaction_data.head())

```

```

      DATE  STORE_NBR  LYLTY_CARD_NBR  TXN_ID  PROD_NBR \
0  43390         1         1000         1         5
1  43599         1         1307         348         66
2  43605         1         1343         383         61
3  43329         2         2373         974         69
4  43330         2         2426        1038        108

      PROD_NAME  PROD_QTY  TOT_SALES \
0  Natural Chip    Compny SeaSalt175g         2         6.0
1      CCs Nacho Cheese    175g         3         6.3
2  Smiths Crinkle Cut Chips Chicken 170g         2         2.9
3  Smiths Chip Thinly S/Cream&Onion 175g         5        15.0
4  Kettle Tortilla ChpsHny&Jlpno Chili 150g         3        13.8

      CLEAN_PROD_WORDS \
0  Natural Chip Compny SeaSaltg
1      CCs Nacho Cheese g
2  Smiths Crinkle Cut Chips Chicken g
3  Smiths Chip Thinly SCreamOnion g
4  Kettle Tortilla ChpsHnyJlpno Chili g

      WORD_FREQUENCY
0      [18645, 6050, 1468, 1468]
1      [246628, 27890, 4658, 4551]
2  [246628, 49770, 28860, 23960, 20754, 15407]
3      [246628, 28860, 18645, 7507, 1473]
4      [246628, 41288, 9580, 3296, 3296]

```

```

In [77]: print(transaction_data.describe())
print(transaction_data.info())
print("check for any null values :\n",transaction_data.isnull().sum())

```

		DATE	STORE_NBR	LYLTY_CARD_NBR	\
count		246742	246742.000000	2.467420e+05	
mean	2018-12-30 01:19:01.211468032		135.051098	1.355310e+05	
min	2018-07-01 00:00:00		1.000000	1.000000e+03	
25%	2018-09-30 00:00:00		70.000000	7.001500e+04	
50%	2018-12-30 00:00:00		130.000000	1.303670e+05	
75%	2019-03-31 00:00:00		203.000000	2.030840e+05	
max	2019-06-30 00:00:00		272.000000	2.373711e+06	
std		NaN	76.787096	8.071528e+04	

	TXN_ID	PROD_NBR	PROD_QTY	TOT_SALES
count	2.467420e+05	246742.000000	246742.000000	246742.000000
mean	1.351311e+05	56.351789	1.908062	7.321322
min	1.000000e+00	1.000000	1.000000	1.700000
25%	6.756925e+04	26.000000	2.000000	5.800000
50%	1.351830e+05	53.000000	2.000000	7.400000
75%	2.026538e+05	87.000000	2.000000	8.800000
max	2.415841e+06	114.000000	200.000000	650.000000
std	7.814772e+04	33.695428	0.659831	3.077828

<class 'pandas.core.frame.DataFrame'>

Index: 246742 entries, 0 to 264835

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	DATE	246742 non-null	datetime64[ns]
1	STORE_NBR	246742 non-null	int64
2	LYLTY_CARD_NBR	246742 non-null	int64
3	TXN_ID	246742 non-null	int64
4	PROD_NBR	246742 non-null	int64
5	PROD_NAME	246742 non-null	object
6	PROD_QTY	246742 non-null	int64
7	TOT_SALES	246742 non-null	float64
8	CLEAN_PROD_WORDS	246742 non-null	object
9	WORD_FREQUENCY	246742 non-null	object

dtypes: datetime64[ns](1), float64(1), int64(5), object(3)

memory usage: 20.7+ MB

None

check for any null values :

DATE	0
STORE_NBR	0
LYLTY_CARD_NBR	0
TXN_ID	0
PROD_NBR	0
PROD_NAME	0
PROD_QTY	0
TOT_SALES	0
CLEAN_PROD_WORDS	0
WORD_FREQUENCY	0

dtype: int64

Filter outliers

```
In [80]: outlier_customer = outliers['LYLTY_CARD_NBR'].unique()[0]

outliers = transaction_data[transaction_data['PROD_QTY'] == 200]
print(outliers)

transaction_data = transaction_data[transaction_data['LYLTY_CARD_NBR'] != outlier_customer]

print(transaction_data.head())
```

	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	PROD_NBR	\
69762	2018-08-19	226	226000	226201	4	
69763	2019-05-20	226	226000	226210	4	

		PROD_NAME	PROD_QTY	TOT_SALES	\
69762	Dorito Corn Chp	Supreme 380g	200	650.0	
69763	Dorito Corn Chp	Supreme 380g	200	650.0	

	CLEAN_PROD_WORDS	WORD_FREQUENCY
69762	Dorito Corn Chp Supreme g	[246628, 22063, 10963, 3185, 3185]
69763	Dorito Corn Chp Supreme g	[246628, 22063, 10963, 3185, 3185]

	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	PROD_NBR	\
0	2018-10-17	1	1000	1	5	
1	2019-05-14	1	1307	348	66	
2	2019-05-20	1	1343	383	61	
3	2018-08-17	2	2373	974	69	
4	2018-08-18	2	2426	1038	108	

		PROD_NAME	PROD_QTY	TOT_SALES	\
0	Natural Chip	Compny SeaSalt175g	2	6.0	
1		CCs Nacho Cheese 175g	3	6.3	
2	Smiths Crinkle Cut Chips Chicken	170g	2	2.9	
3	Smiths Chip Thinly S/Cream&Onion	175g	5	15.0	
4	Kettle Tortilla ChpsHny&Jlpno Chili	150g	3	13.8	

	CLEAN_PROD_WORDS	\
0	Natural Chip Compny SeaSaltg	
1	CCs Nacho Cheese g	
2	Smiths Crinkle Cut Chips Chicken g	
3	Smiths Chip Thinly SCreamOnion g	
4	Kettle Tortilla ChpsHnyJlpno Chili g	

	WORD_FREQUENCY
0	[18645, 6050, 1468, 1468]
1	[246628, 27890, 4658, 4551]
2	[246628, 49770, 28860, 23960, 20754, 15407]
3	[246628, 28860, 18645, 7507, 1473]
4	[246628, 41288, 9580, 3296, 3296]

```
In [83]: transactions_counts = transaction_data.groupby('DATE').size().reset_index(name='TRANSACTION_COUNT')
print(transactions_by_date.head())
transactions_by_date.count()
```

```
Out[83]:
```

	DATE	TRANSACTION_COUNT
0	2018-07-01	663
1	2018-07-02	650
2	2018-07-03	674
3	2018-07-04	669
4	2018-07-05	660

```
DATE 364
TRANSACTION_COUNT 364
dtype: int64
```

As we see in the above data, it is shown that there are only 364 rows, meaning only 364 dates only which indicates two dates are missing.

A chart of number of transactions over time to find the missing date.

```
In [89]: full_date_range = pd.date_range(start='2018-07-01', end='2019-06-30')

transaction_counts = transaction_data.groupby('DATE').size().reset_index(name='TRANSACTION_COUNT')

transaction_data_daily = pd.DataFrame({'DATE': full_date_range})

transaction_data_daily = transaction_data_daily.merge(transaction_counts, on='DATE', how='left')

transaction_data_daily['TRANSACTION_COUNT'] = transaction_data_daily['TRANSACTION_COUNT'].fillna(0)

missing_dates = transaction_data_daily[transaction_data_daily['TRANSACTION_COUNT'] == 0]['DATE']
print("Missing date(s):")
print(missing_dates.dt.strftime('%Y-%m-%d').tolist())

Missing date(s):
['2018-12-25']
```

```
In [90]: plt.figure(figsize=(14, 6))
sns.lineplot(data=transaction_data_daily, x='DATE', y='TRANSACTION_COUNT')


for date in missing_dates:
    plt.axvline(x=date, color='red', linestyle='--', alpha=0.7, label='Missing Date')

plt.title('Transactions Over Time')
plt.xlabel('Date')
plt.ylabel('Number of Transactions')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

```

/usr/local/lib/python3.11/dist-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
/usr/local/lib/python3.11/dist-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):

```



```

In [91]: december_data = transaction_data[transaction_data['DATE'].dt.month == 12]

december_by_day = december_data.groupby('DATE').size().reset_index(name='TRANSACTION_COUNT')

print(december_by_day.head())

```

	DATE	TRANSACTION_COUNT
0	2018-12-01	675
1	2018-12-02	655
2	2018-12-03	677
3	2018-12-04	666
4	2018-12-05	660

```


In [92]: plt.figure(figsize=(12, 6))
sns.lineplot(data=december_by_day, x='DATE', y='TRANSACTION_COUNT')
plt.title('Number of Transactions in December (by Day)')
plt.xlabel('Date')
plt.ylabel('Transaction Count')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```

```

/usr/local/lib/python3.11/dist-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
/usr/local/lib/python3.11/dist-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):

```



From this it is clear that the increase in the sale in lead-up to christmas and that too zero sales on the christmas day

itself. This is mainly due to shops being closed on christmas day.


```
In [94]: transaction_data['PACK_SIZE'] = transaction_data['PROD_NAME'].str.extract(r'(\d+)\s*[gG]')[0].astype(float)

pack_size_summary = transaction_data['PACK_SIZE'].value_counts().sort_index()
print(pack_size_summary)
```

```
PACK_SIZE
70.0      1507
90.0      3008
110.0     22387
125.0     1454
134.0     25102
135.0      3257
150.0     40203
160.0      2970
165.0     15297
170.0     19983
175.0     66390
180.0      1468
190.0      2995
200.0      4473
210.0      6272
220.0      1564
250.0      3169
270.0      6285
330.0     12540
380.0      6416
Name: count, dtype: int64
```

```
In [95]: sns.set(style="whitegrid")

plt.figure(figsize=(10, 6))
sns.countplot(data=transaction_data, x='PACK_SIZE', order=sorted(transaction_data['PACK_SIZE'].dropna().unique()))
plt.title('Number of Transactions by Pack Size')
plt.xlabel('Pack Size (g)')
plt.ylabel('Number of Transactions')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
In [96]: transaction_data['BRAND'] = transaction_data['PROD_NAME'].str.split().str[0]
```

```
In [97]: transaction_data['BRAND'] = transaction_data['BRAND'].replace({
    'RED': 'RRD',
    'RRD': 'RRD',
    'Infzns': 'Infuzions',
    'Infzn': 'Infuzions',
    'Snbts': 'Sunbites',
    'Dorito': 'Doritos',
    'WW': 'Woolworths',
    'Smith': 'Smiths'
    # You can add more replacements here if needed
})
```

```
In [98]: print(transaction_data[['PROD_NAME', 'BRAND']].head(10))
print("\nBrand counts:\n")
print(transaction_data['BRAND'].value_counts())
```

```

          PROD_NAME  BRAND
0  Natural Chip      Compny SeaSalt175g  Natural
1          CCs Nacho Cheese    175g      CCs
2  Smiths Crinkle Cut  Chips Chicken 170g  Smiths
3  Smiths Chip Thinly  S/Cream&Onion 175g  Smiths
4  Kettle Tortilla ChpsHny&Jlpno Chili 150g  Kettle
6  Smiths Crinkle Chips Salt & Vinegar 330g  Smiths
7    Grain Waves      Sweet Chilli 210g   Grain
8  Doritos Corn Chip Mexican Jalapeno 150g  Doritos
9    Grain Waves Sour    Cream&Chives 210G   Grain
10 Smiths Crinkle Chips Salt & Vinegar 330g  Smiths

```

Brand counts:

```

BRAND
Kettle      41288
Smiths      30353
Doritos     25224
Pringles    25102
Infuzions   14201
Thins       14075
RRD         11894
Woolworths  11836
Cobs        9693
Tostitos    9471
Twisties    9454
Tyrrells    6442
Grain       6272
Natural     6050
Cheezels    4603
CCs         4551
Red         4427
Sunbites    3008
Cheetos     2927
Burger      1564
GrnWves     1468
NCC         1419
French      1418
Name: count, dtype: int64

```

```
In [102]: customer_data.info()

print("LIFESTAGE Distribution:\n")
print(customer_data['LIFESTAGE'].value_counts())

print("\nPREMIUM_CUSTOMER Distribution:\n")
print(customer_data['PREMIUM_CUSTOMER'].value_counts())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 72637 entries, 0 to 72636
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   LYLTY_CARD_NBR   72637 non-null  int64
1   LIFESTAGE        72637 non-null  object
2   PREMIUM_CUSTOMER 72637 non-null  object

```

dtypes: int64(1), object(2)

memory usage: 1.7+ MB

LIFESTAGE Distribution:

```

LIFESTAGE
RETIREEES      14805
OLDER SINGLES/COUPLES 14609
YOUNG SINGLES/COUPLES 14441
OLDER FAMILIES    9780
YOUNG FAMILIES    9178
MIDAGE SINGLES/COUPLES 7275
NEW FAMILIES      2549
Name: count, dtype: int64

```

PREMIUM_CUSTOMER Distribution:

```

PREMIUM_CUSTOMER
Mainstream    29245
Budget        24470
Premium       18922
Name: count, dtype: int64

```

```
In [104]: data = transaction_data.merge(customer_data, on='LYLTY_CARD_NBR', how='left')
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 246740 entries, 0 to 246739
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   DATE                  246740 non-null  datetime64[ns]
1   STORE_NBR             246740 non-null  int64
2   LYLTY_CARD_NBR        246740 non-null  int64
3   TXN_ID                246740 non-null  int64
4   PROD_NBR              246740 non-null  int64
5   PROD_NAME             246740 non-null  object
6   PROD_QTY              246740 non-null  int64
7   TOT_SALES             246740 non-null  float64
8   CLEAN_PROD_WORDS      246740 non-null  object
9   WORD_FREQUENCY        246740 non-null  object
10  PACK_SIZE             246740 non-null  float64
11  BRAND                 246740 non-null  object
12  LIFESTAGE             246740 non-null  object
13  PREMIUM_CUSTOMER      246740 non-null  object
dtypes: datetime64[ns](1), float64(2), int64(5), object(6)
memory usage: 26.4+ MB
```

```
In [108.] print(data.isnull().sum())
print(data[data[['LIFESTAGE', 'PREMIUM_CUSTOMER']].isnull().any(axis=1)])

DATE                0
STORE_NBR           0
LYLTY_CARD_NBR      0
TXN_ID              0
PROD_NBR            0
PROD_NAME           0
PROD_QTY            0
TOT_SALES           0
CLEAN_PROD_WORDS    0
WORD_FREQUENCY      0
PACK_SIZE           0
BRAND               0
LIFESTAGE           0
PREMIUM_CUSTOMER    0
dtype: int64
Empty DataFrame
Columns: [DATE, STORE_NBR, LYLTY_CARD_NBR, TXN_ID, PROD_NBR, PROD_NAME, PROD_QTY, TOT_SALES, CLEAN_PROD_WORDS, WORD_FREQUENCY, PACK_SIZE, BRAND, LIFESTAGE, PREMIUM_CUSTOMER]
Index: []
```

```
In [109.] data.to_csv("QVI_data.csv", index=False)
```

- Who spends the most on chips (total sales), describing customers by lifestage and how premium their general purchasing behaviour is
- How many customers are in each segment
- How many chips are bought per customer by segment
- What's the average chip price by customer segment

```
In [110.] total_sales_by_segment = data.groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER'])['TOT_SALES'].sum().reset_index()
print(total_sales_by_segment)
```

	LIFESTAGE	PREMIUM_CUSTOMER	TOT_SALES
0	MIDAGE SINGLES/COUPLES	Budget	33345.70
1	MIDAGE SINGLES/COUPLES	Mainstream	84734.25
2	MIDAGE SINGLES/COUPLES	Premium	54443.85
3	NEW FAMILIES	Budget	20607.45
4	NEW FAMILIES	Mainstream	15979.70
5	NEW FAMILIES	Premium	10760.80
6	OLDER FAMILIES	Budget	156863.75
7	OLDER FAMILIES	Mainstream	96413.55
8	OLDER FAMILIES	Premium	75242.60
9	OLDER SINGLES/COUPLES	Budget	127833.60
10	OLDER SINGLES/COUPLES	Mainstream	124648.50
11	OLDER SINGLES/COUPLES	Premium	123537.55
12	RETIREEES	Budget	105916.30
13	RETIREEES	Mainstream	145168.95
14	RETIREEES	Premium	91296.65
15	YOUNG FAMILIES	Budget	129717.95
16	YOUNG FAMILIES	Mainstream	86338.25
17	YOUNG FAMILIES	Premium	78571.70
18	YOUNG SINGLES/COUPLES	Budget	57122.10
19	YOUNG SINGLES/COUPLES	Mainstream	147582.20
20	YOUNG SINGLES/COUPLES	Premium	39052.30

```
In [111.] # Plotting total sales by LIFESTAGE and PREMIUM_CUSTOMER
plt.figure(figsize=(10, 6))
sns.barplot(data=total_sales_by_segment, x='LIFESTAGE', y='TOT_SALES', hue='PREMIUM_CUSTOMER')

plt.xlabel('Lifestage')
plt.ylabel('Total Sales')
```

```
plt.title('Total Sales by Lifestage and Premium Customer Segment')
plt.xticks(rotation=45)
plt.tight_layout()
plt.legend(title='Premium Customer')
plt.show()
```




Sales are coming mainly from Budget - older families, Mainstream - retirees, and Mainstream - young single / couples.

```
In [112]: num_customers_by_segment = data.groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER'])['LYLTY_CARD_NBR'].nunique().reset_index()
print(num_customers_by_segment)
```

	LIFESTAGE	PREMIUM_CUSTOMER	LYLTY_CARD_NBR
0	MIDAGE SINGLES/COUPLES	Budget	1474
1	MIDAGE SINGLES/COUPLES	Mainstream	3298
2	MIDAGE SINGLES/COUPLES	Premium	2369
3	NEW FAMILIES	Budget	1087
4	NEW FAMILIES	Mainstream	830
5	NEW FAMILIES	Premium	575
6	OLDER FAMILIES	Budget	4611
7	OLDER FAMILIES	Mainstream	2788
8	OLDER FAMILIES	Premium	2231
9	OLDER SINGLES/COUPLES	Budget	4849
10	OLDER SINGLES/COUPLES	Mainstream	4858
11	OLDER SINGLES/COUPLES	Premium	4682
12	RETIREES	Budget	4385
13	RETIREES	Mainstream	6358
14	RETIREES	Premium	3812
15	YOUNG FAMILIES	Budget	3953
16	YOUNG FAMILIES	Mainstream	2685
17	YOUNG FAMILIES	Premium	2398
18	YOUNG SINGLES/COUPLES	Budget	3647
19	YOUNG SINGLES/COUPLES	Mainstream	7917
20	YOUNG SINGLES/COUPLES	Premium	2480

```
In [113]: # Plotting the number of customers by LIFESTAGE and PREMIUM_CUSTOMER
plt.figure(figsize=(10, 6))
sns.barplot(data=num_customers_by_segment, x='LIFESTAGE', y='LYLTY_CARD_NBR', hue='PREMIUM_CUSTOMER')

plt.xlabel('Lifestage')
plt.ylabel('Number of Customers')
plt.title('Number of Customers by Lifestage and Premium Customer Segment')
plt.xticks(rotation=45)
plt.tight_layout()
plt.legend(title='Premium Customer')
plt.show()
```



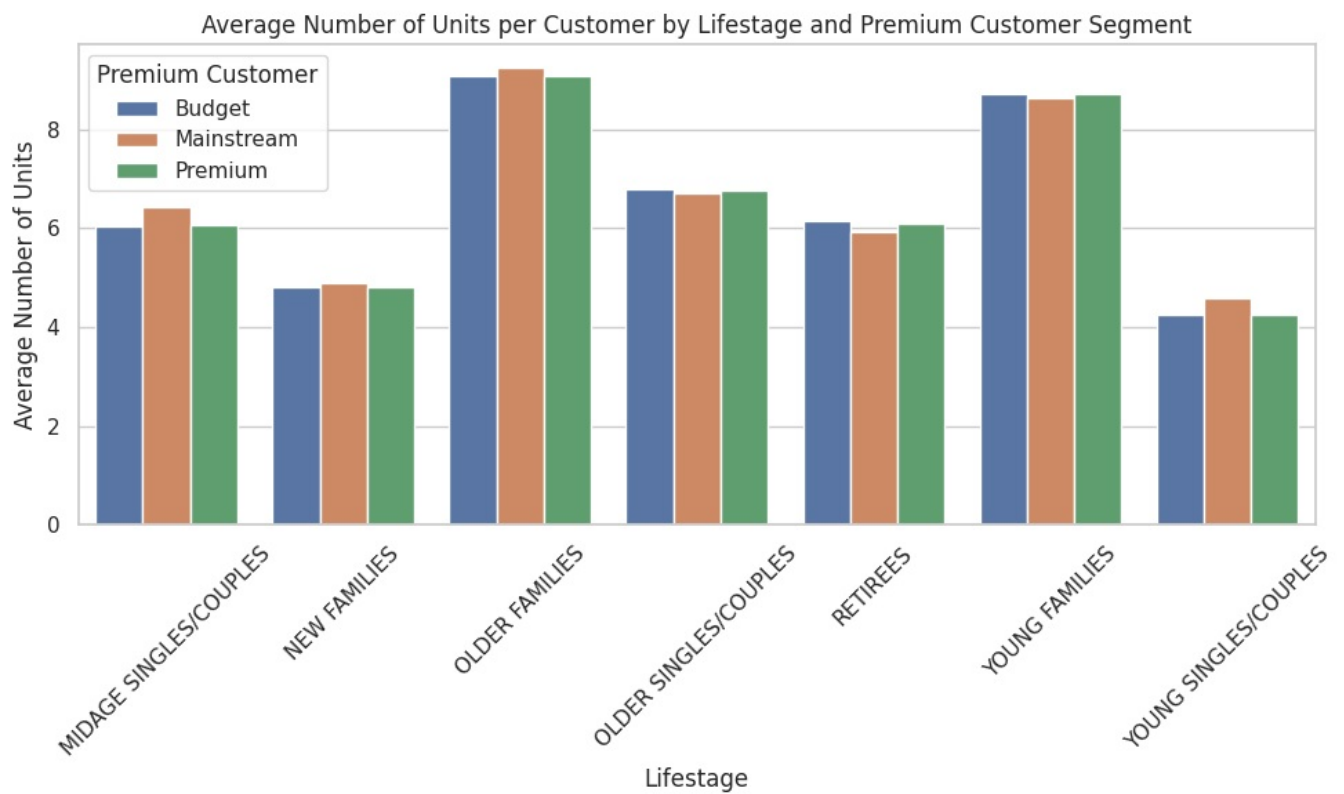
There are more Mainstream - young singles/couples and Mainstream - retirees who buy chips. This contributes to there being more sales to these customer segments but this is not a major driver for the Budget - Older families segment.

```
In [114]: total_units_per_customer = data.groupby(['LYLTY_CARD_NBR', 'LIFESTAGE', 'PREMIUM_CUSTOMER'])['PROD_QTY'].sum()
avg_units_per_customer = total_units_per_customer.groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER'])['PROD_QTY'].mean()
print(avg_units_per_customer)
```

	LIFESTAGE	PREMIUM_CUSTOMER	PROD_QTY
0	MIDAGE SINGLES/COUPLES	Budget	6.026459
1	MIDAGE SINGLES/COUPLES	Mainstream	6.432080
2	MIDAGE SINGLES/COUPLES	Premium	6.078514
3	NEW FAMILIES	Budget	4.821527
4	NEW FAMILIES	Mainstream	4.891566
5	NEW FAMILIES	Premium	4.815652
6	OLDER FAMILIES	Budget	9.076773
7	OLDER FAMILIES	Mainstream	9.255380
8	OLDER FAMILIES	Premium	9.071717
9	OLDER SINGLES/COUPLES	Budget	6.781398
10	OLDER SINGLES/COUPLES	Mainstream	6.712021
11	OLDER SINGLES/COUPLES	Premium	6.769543
12	RETIREES	Budget	6.141847
13	RETIREES	Mainstream	5.925920
14	RETIREES	Premium	6.103358
15	YOUNG FAMILIES	Budget	8.722995
16	YOUNG FAMILIES	Mainstream	8.638361
17	YOUNG FAMILIES	Premium	8.716013
18	YOUNG SINGLES/COUPLES	Budget	4.250069
19	YOUNG SINGLES/COUPLES	Mainstream	4.575597
20	YOUNG SINGLES/COUPLES	Premium	4.264113

```
In [115]: # Plotting the average number of units per customer by LIFESTAGE and PREMIUM_CUSTOMER
plt.figure(figsize=(10, 6))
sns.barplot(data=avg_units_per_customer, x='LIFESTAGE', y='PROD_QTY', hue='PREMIUM_CUSTOMER')

plt.xlabel('Lifestage')
plt.ylabel('Average Number of Units')
plt.title('Average Number of Units per Customer by Lifestage and Premium Customer Segment')
plt.xticks(rotation=45)
plt.tight_layout()
plt.legend(title='Premium Customer')
plt.show()
```



Older families and young families in general buy more chips per customer.

```
In [116.. data['PRICE_PER_UNIT'] = data['TOT_SALES'] / data['PROD_QTY']
print(data[['LYLTY_CARD_NBR', 'PROD_QTY', 'TOT_SALES', 'PRICE_PER_UNIT']].head())
```


	LYLTY_CARD_NBR	PROD_QTY	TOT_SALES	PRICE_PER_UNIT
0	1000	2	6.0	3.00
1	1307	3	6.3	2.10
2	1343	2	2.9	1.45
3	2373	5	15.0	3.00
4	2426	3	13.8	4.60

```
In [117.. avg_price_per_unit = data.groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER'])['PRICE_PER_UNIT'].mean().reset_index()
print(avg_price_per_unit)
```

	LIFESTAGE	PREMIUM_CUSTOMER	PRICE_PER_UNIT
0	MIDAGE SINGLES/COUPLES	Budget	3.743328
1	MIDAGE SINGLES/COUPLES	Mainstream	3.994241
2	MIDAGE SINGLES/COUPLES	Premium	3.770698
3	NEW FAMILIES	Budget	3.917688
4	NEW FAMILIES	Mainstream	3.916133
5	NEW FAMILIES	Premium	3.872110
6	OLDER FAMILIES	Budget	3.745340
7	OLDER FAMILIES	Mainstream	3.737077
8	OLDER FAMILIES	Premium	3.717000
9	OLDER SINGLES/COUPLES	Budget	3.882096
10	OLDER SINGLES/COUPLES	Mainstream	3.814665
11	OLDER SINGLES/COUPLES	Premium	3.893182
12	RETIREES	Budget	3.924404
13	RETIREES	Mainstream	3.844294
14	RETIREES	Premium	3.920942
15	YOUNG FAMILIES	Budget	3.760737
16	YOUNG FAMILIES	Mainstream	3.724533
17	YOUNG FAMILIES	Premium	3.762150
18	YOUNG SINGLES/COUPLES	Budget	3.657366
19	YOUNG SINGLES/COUPLES	Mainstream	4.065642
20	YOUNG SINGLES/COUPLES	Premium	3.665414

```
In [118.. # Plotting the average price per unit by LIFESTAGE and PREMIUM_CUSTOMER
plt.figure(figsize=(10, 6))
sns.barplot(data=avg_price_per_unit, x='LIFESTAGE', y='PRICE_PER_UNIT', hue='PREMIUM_CUSTOMER')

plt.xlabel('Lifestage')
plt.ylabel('Average Price per Unit')
plt.title('Average Price per Unit by Lifestage and Premium Customer Segment')
plt.xticks(rotation=45)
plt.tight_layout()
plt.legend(title='Premium Customer')
plt.show()
```



Mainstream midage and young singles and couples are more willing to pay more per packet of chips compared to their budget and premium counterparts. This may be due to premium shoppers being more likely to buy healthy snacks and when they buy chips, this is mainly for entertainment purposes rather than their own consumption. This is also supported by there being fewer premium midage and young singles and couples buying chips compared to their mainstream counterparts.

```
In [119... from scipy.stats import ttest_ind

mainstream_data = data[data['PREMIUM_CUSTOMER'] == 'Mainstream']['PRICE_PER_UNIT']
premium_data = data[data['PREMIUM_CUSTOMER'] == 'Premium']['PRICE_PER_UNIT']

t_stat, p_value = ttest_ind(mainstream_data, premium_data, equal_var=False) # Welch's t-test

print(f"T-statistic: {t_stat}")
print(f"P-value: {p_value}")

T-statistic: 11.05723574336515
P-value: 2.078836404116925e-28
```

In []: